



CentraleSupélec



From Multistage Stochastic Programming to Markov Decision Process

Yiping Fang, Professor

Chair on Risk and Resilience of Complex Systems,

Laboratoire Génie Industriel, CentraleSupélec, Université Paris Saclay

Yiping.fang@centralesupelec.fr




- Multistage stochastic programming
- A different perspective: Markov decision processes
 - modeling elements
 - solving MDP: value iteration and policy iteration
 - dual perspective: linear programming for MDP



- We take decisions in two-stages

$$x_0 \rightsquigarrow \xi_1 \rightsquigarrow x_1$$



- A scenario tree is a **graphical representation** of a stochastic decision process
- On a scenario tree, it means solving the **extensive formulation** (linear):

$$\min_{x_0, x_{1,s}} \quad c_0^\top x_0 + \sum_{s \in \mathcal{S}} p_s [c_s^\top x_{1,s}]$$


$$\text{s.t.} \quad x_0 \in \mathcal{X}_0, x_{1,s} \in \mathcal{X}_1(x_0, \xi_{1,s})$$



Generalize to multistage



CentraleSupélec




- Ubiquitous
- Highly challenging, modeling and computationally
- Exogenous uncertainty: uncertainty realizations are not influenced by decisions



Generalize to multistage




CentraleSupélec




- We take decisions in T stages
 $x_0 \rightsquigarrow \xi_1 \rightsquigarrow x_1 \rightsquigarrow \dots \rightsquigarrow \xi_T \rightsquigarrow x_T$
- Can be represented on a tree $\mathcal{T}$, where a node n of depth t represent a realization of
 $(\xi_1, \dots, \xi_t) : p_n$
- Then, the **extensive form** where we **minimize expect total cost** reads

$$\min_{\{x_n\}_{n \in \mathcal{T}}} \sum_{n \in \mathcal{T}} p_n c_n(x_n)$$



- Consider a tree of depth T . A scenario $\{s\} := (n_1, n_2, \dots, n_T)$ uniquely defined by its last element (leaf) $p_s = p_{n_T}$
- The splitted scenario formulation reads

$$\begin{aligned} \min_{\{x_n\}_{n \in \mathcal{T}}} \quad & \sum_{n \in \mathcal{T}} p_n c_n(x_n) \\ &= \sum_{s \in S} p_s \sum_{n \in N_s} c_n(x_n) \end{aligned}$$



- The **splitted scenario** formulation

$$\begin{aligned} \min_{\{x_{t,s}\}_{s \in S, t \in \{0, \dots, T\}}} \quad & \sum_{s \in S} p_s \sum_{t=0}^T c_{t,s}(x_{t,s}) \\ \text{s.t.} \quad & x_{t,s} = x_{t,s'}, \\ & \forall t, n \in \mathcal{N}_t, \{s\} \cap \{s'\} \ni n \end{aligned}$$

where $\mathcal{N}_t$ is the set of nodes of depth t

- Nonanticipativity**: at time t , decisions are taken only knowing the past realizations of the uncertainty




Example: hydro power planning



CentraleSupélec

- How much hydro power to generate in each period to satisfy demand with minimum expected cost?



$$\min_{p_t, q_t} \sum_t c_t^e q_t + c^p p_t$$

$$\text{s.t. } h_t = h_{t-1} + \xi_t - p_{t-1}, \forall t$$

$$\alpha p_t + q_t = d_t, \forall t$$

$$0 \leq p_t \leq h_t \leq h^{\max}, \forall t$$

$$p_t, q_t \geq 0, \forall t$$

- Generation cost: $c^p p_t$
- Buying remote electricity cost: $c_t^e q_t$
- Stochastic inflow ξ_t**





Example: hydro power planning



CentraleSupélec

- Stochastic inflows ξ_t follow the scenario tree
- Extensive form?
- Splitted scenario form?



$$p_{s=2} = p_{n_4} = 0.8 \cdot 0.3 = 0.24$$




Example: hydro power planning



CentraleSupélec

- Extensive form



Node	Label	pa	p_n
n_0	root	$\emptyset$	1.00
n_1	$\xi_{1,1}$	n_0	0.80
n_2	$\xi_{1,2}$	n_0	0.20
n_3	$\xi_{1,1} \rightarrow \xi_{2,1}$	n_1	$0.80 \cdot 0.7 = 0.56$
n_4	$\xi_{1,1} \rightarrow \xi_{2,2}$	n_1	$0.80 \cdot 0.3 = 0.24$
n_5	$\xi_{1,2} \rightarrow \xi_{2,1}$	n_2	$0.20 \cdot 0.7 = 0.14$
n_6	$\xi_{1,2} \rightarrow \xi_{2,2}$	n_2	$0.20 \cdot 0.3 = 0.06$

$$\begin{aligned}
 & \min_{\{p_n, q_n, h_n\}_{n \in \mathcal{T}}} \sum_{n \in \mathcal{T}} p_n (c_n^e q_n + c^p p_n) \\
 & \text{s.t.} \quad h_n = h_{\text{pa}(n)} + \xi_n - p_{\text{pa}(n)}, \forall n \neq n_0 \\
 & \quad \alpha p_n + q_n = d_n, \forall n \in \mathcal{T} \\
 & \quad 0 \leq p_n \leq h_n \leq h^{\max}, \forall n \in \mathcal{T} \\
 & \quad p_n, q_n \geq 0, \forall n \in \mathcal{T}
 \end{aligned}$$




Example: hydro power planning



CentraleSupélec

- Splitted scenario form

Scenario s	Path of (ξ_1, ξ_2)	Probability p_s
$s^{(1)}$	$(\xi_{1,1}, \xi_{2,1})$	$0.8 \cdot 0.7 = 0.56$
$s^{(2)}$	$(\xi_{1,1}, \xi_{2,2})$	$0.8 \cdot 0.3 = 0.24$
$s^{(3)}$	$(\xi_{1,2}, \xi_{2,1})$	$0.2 \cdot 0.7 = 0.14$
$s^{(4)}$	$(\xi_{1,2}, \xi_{2,2})$	$0.2 \cdot 0.3 = 0.06$





- The splitted scenario formulation (although LP) is usually **too large** to be solved as one monolithic problem.
- **Decomposition:**
 - Stage-wise decomposition: nested Benders
 - Stochastic Dual Dynamic Programming (state-of-the-art for large-scale multistage)
- The **complexity grows exponentially** with the number of depth (stage) T




- Multistage stochastic programming
- A different perspective: Markov decision processes
 - modeling elements
 - solving MDP: value iteration and policy iteration
 - dual perspective: linear programming for MDP




History and connection



CentraleSupélec



• ... seminal paper ○ ... chapter in ⊗ selected monograph



Richard Bellman

On the theory of
dynamic programming
(1952)

A Markovian Decision
Process (1957)



Ronald A. Howard

Dynamic programming
and Markov Processes,
1960 (Book)

MDP: Markov decision processes

DP: Dynamic programming

(M)SP: (Multistage) Stochastic
programming

ChC: Chance-constraint




- **Motivations** for MDP:
 - very long horizon
 - interested in finding a **decision rule/policy** that depends on the **observed system state**
- An appropriate **definition of system state** is the central point in MDP; whereas states **not necessarily appear explicitly** in traditional MSP



State and policy in MSP



CentraleSupélec



$$\mathbf{s}_t = (h_t, d_t, c_t^e)$$

$$\mathbf{a}_t = (p_t, q_t)$$

$$\mathbf{s}_{t+1} = (h_{t+1} = h_t - p_t + \xi_{t+1}, d_{t+1}, c_{t+1}^e)$$

$$\mathbf{a}_{t+1} = (p_{t+1}, q_{t+1})$$

$$\min_{p_t, q_t} \quad c_t^e q_t + c^p p_t + \mathcal{Q}(\mathbf{s}_t, \mathbf{a}_t)$$

$$\begin{aligned} \text{s.t.} \quad & \alpha p_t + q_t = d_t \\ & 0 \leq p_t \leq h_t \leq h^{max}, \\ & p_t, q_t \geq 0 \end{aligned}$$


$$\min_{p_{t+1}, q_{t+1}} \quad c_{t+1}^e q_{t+1} + c^p p_{t+1} + \mathcal{Q}(\mathbf{s}_{t+1}, \mathbf{a}_{t+1})$$

$$\begin{aligned} \text{s.t.} \quad & \alpha p_{t+1} + q_{t+1} = d_{t+1} \\ & 0 \leq p_{t+1} \leq h_{t+1} \leq h^{max}, \\ & p_{t+1}, q_{t+1} \geq 0 \end{aligned}$$

$$\text{Policy: } (h_t, d_t, c_t^e) \xrightarrow{\pi(\mathbf{a}_t | \mathbf{s}_t)} (p_t, q_t)$$



- Exogenous vs. endogenous uncertainty
 - **Exogenous**, e.g., inflows
 - **Endogenous**, e.g., degradation rate of a robot car depends on the driving behaviors (fast, slow, etc.) $\Rightarrow$ not easy modeled in MSP



- MDP is a natural fit in this case, but at the sacrifice of assuming **discrete (small) decision space** (traditionally)





Typical MDP example: racing



CentraleSupélec

- A robot car wants to travel far, quickly
- 3 **states**: *Cool*, *Warm*, Overheated
- 2 **actions**: *Slow*, *Fast*
- Going faster gets double **reward** as long as not overheated
- **Goal**: maximize sum of rewards






Example: grid world




CentraleSupélec

- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- **Noisy movement:** actions do not always go as planned
 - $p = 0.8$, the action **North** takes the agent North (if no wall there)
 - $p = 0.1$, **North** takes the agent West; $p = 0.1$ East
 - If there is a wall in the direction, the agent stays put
- The agent **receives rewards** each time step
 - Small “living” reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- **Goal:** maximize sum of rewards





- A **MDP** is defined by
 - a set of **system states** $s \in S$
 - a set of **actions/decisions**, $a \in A$
 - a **transition function** $T(s, a, s')$
 - probability that action a at state s leads to s' , i.e. $P(s'|s, a)$
 - also called the dynamics, **natural fit of endogenous uncertainty**
 - a **reward function** $R(s, a, s')$
 - sometimes just $R(a)$ or $R(s, a)$
 - initial state is given






- **Markov**: given the present state, the future and the past are independent
- For MDPs, “Markov” means **action outcomes depends only on the current state & action**

$$\begin{aligned} &P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ &= P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$



- Each MDP state projects a scenario search tree

chance node,
Q-state (s, a) :
e.g., drive fast
when in warm
state







Racing search tree



CentraleSupélec



- In MSP, we want an optimal sequence of decisions, **one for each node in the scenario tree**, $\Xi \rightarrow A$
- For MDPs, we want an **optimal policy** $\pi^* : S \rightarrow A$
 - a policy $\pi : S \rightarrow A$ gives an action for each state
 - **optimal policy**: **maximizes expected utility** if followed
 - an explicit policy defines a **reflex agent**: reduce to a Markov reward process






Utility of sequence



CentraleSupélec

- What preferences should an agent have over reward sequences?
- More or less?
 $[1, 2, 2]$ or $[2, 3, 4]$
- Now or later?
 $[1, 0, 0]$ or $[0, 0, 1]$





- It's reasonable to maximize **the sum of rewards** of the sequence
- It's also reasonable to **prefer rewards now** to rewards later
- **One solution:** values of rewards decay exponentially



1

Worth now



γ

Worth next step



γ^2


Worth in 2 steps

$$U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$$

- **Example:** $\gamma = 0.5$
 - $\rightarrow U([1, 2, 3]) = 1 \times 1 + 0.5 \times 2 + 0.25 \times 3$
 - $\rightarrow U([1, 2, 3]) < U([3, 2, 1])$



- What if the horizon is infinite?
→ infinite rewards?



- Solutions
→ use **discounted utility** $0 < \gamma < 1$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq r_{\max} / (1 - \gamma)$$

- smaller γ means smaller “horizon”, **shorter term** focus
- **absorbing state**: guarantee that for every policy, a terminal state will eventually be reached



- A **value function** for a policy π , written $V^\pi : S \rightarrow \mathbf{R}$ gives the **expected utility** (sum of discounted rewards) when acting under that policy

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \mid s_0 = s, a_t = \pi(s_t) \right]$$

- The **optimal policy** is the policy that achieves the highest value for every state

$$\pi^\star = \arg \max_{\pi} V^\pi(s)$$


- and its **optimal value** is written $V^\star(s)$




Optimal policy




CentraleSupélec




$$R(s) = -0.01$$



$$R(s) = -0.03$$



$$R(s) = -0.4$$



$$R(s) = -2.0$$



Recap: defining MDPs



CentraleSupélec

- Markov decision processes
 - System states $s \in S$
 - Start state s_0
 - Possible actions $a \in A$
 - Transitions $T(s, a, s') = P(s'|s, a)$
 - Rewards $R(s, a, s')$ and discount γ
- MDP quantities so far
 - utility: **sum of (discounted) rewards**
 - policy: **choice of action for each state**
 - value function for a given policy: **expected utility** when acting under that policy
 - optimal policy: gives the **highest value** for every state




Tutorial: parking problem



CentraleSupélec

- You are driving to a restaurant in Paris
- start from street parking spot N , and you can observe whether parking spot $n, n \in \{N, N-1, \dots, 0\}$ is empty **only when arriving at the spot**.
- by p , a spot is empty; $1-p$, a spot is occupied.
- the parking lot at the restaurant is always available with a high fee c .
- The cost of parking at a street spot n is equal to the distance from the restaurant n





MDP formulation

- States

$$s \in S \cup \{(0, 1), \Delta\}, \text{ where } S = \mathbf{Z}_+ \times \{0, 1\}$$

- Actions

$$A_s = \begin{cases} \{C\} & s = (n, 0) \\ \{C, P\} & s = (n, 1) \\ \{P\} & s = (0, 1) \end{cases}$$

- Costs (rewards)

$$R_n(s, a) = \begin{cases} 0 & s \in S, a = C \\ n & s \in S, a = P \\ c & s = (0, 1), a = P \end{cases}$$



MDP formulation


- Transition probabilities

$$T((n, \cdot), C, (n-1, 1)) = p, \forall n \geq 2$$

$$T((n, \cdot), C, (n-1, 0)) = 1 - p, \forall n \geq 2$$

$$T((1, \cdot), C, (0, 1)) = 1$$

$$T((n, 1), P, \Delta) = 1$$





- Indicating parameter


$$\delta_{n,s} = \begin{cases} 0 & \text{if spot } n \text{ is occupied in scenario } s \\ 1 & \text{if spot } n \text{ is empty in scenario } s \end{cases}$$

- Binary decision variable

$$y_{n,s} = \begin{cases} 0 & \text{if do not park in spot } n \text{ in scenario } s \\ 1 & \text{if park in spot } n \text{ in scenario } s \end{cases}$$

- Probability of scenario s

$$p_s = \prod_{n=1}^N (p\delta_{n,s} + (1-p)(1-\delta_{n,s}))$$





- MILP formulation

$$\min_{y_{n,s}} \sum_s p_s c_s$$

$$\text{s.t. } c_s = \sum_{n=1}^N y_{n,s} n + (1 - \sum_{n=1}^N y_{n,s}) c, \quad \forall s \in S$$

$$y_{n,s} \leq \delta_{n,s}, \quad \forall n = 1, \dots, N, s \in S$$

$$\sum_{n=1}^N y_{n,s} \leq 1, \quad \forall s \in S$$

$$y_{n,s} = y_{n,s'}, \quad \forall N_{s,s'}^{\text{parent}} \leq n \leq N, s, s' \in S$$



- For example: $N = 3, p = 0.6, c = 4$

$$\delta_{n,s} = \begin{cases} 0 & \text{if spot } n \text{ is occupied in scenario } s \\ 1 & \text{if spot } n \text{ is empty in scenario } s \end{cases}$$

$$y_{n,s} = \begin{cases} 0 & \text{if do not park in spot } n \text{ in scenario } s \\ 1 & \text{if park in spot } n \text{ in scenario } s \end{cases}$$


	$s = 1$		$s = 2$		$s = 3$		$s = 4$		$s = 5$		$s = 6$		$s = 7$		$s = 8$	
n	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$	$\delta_{n,s}$	$y_{n,s}$
1	0	0	1	1	0	0	1	0	0	0	1	1	0	0	1	0
2	0	0	0	0	1	1	1	1	0	0	0	0	1	1	1	1
3	0	0	0	0	0	0	0	0	1	0	1	0	1	0	1	0



$$V[(n, 0)] = pV[(n - 1, 1)] + (1 - p)V[(n - 1, 0)]$$

$$V[(n, 1)] = \min \left\{ \underbrace{n}_{\text{To park}}, \underbrace{pV[(n - 1, 1)] + (1 - p)V[(n - 1, 0)]}_{\text{To keep looking}} \right\}$$

$$N = 3, p = 0.6, c = 4$$



$$\begin{aligned} V[(3, 0)] &= 0.6V[(2, 1)] + 0.4V[(2, 0)] \\ &= 0.6 \times 2 + 0.4 \times 2.2 = 2.08 \end{aligned}$$

$$V[(3, 1)] = \min\{3, 2.08\} = 2.08$$

To keep looking

$$\begin{aligned} V[(2, 0)] &= 0.6V[(1, 1)] + 0.4V[(1, 0)] \\ &= 0.6 \times 1 + 0.4 \times 4 = 2.2 \end{aligned}$$

$$V[(2, 1)] = \min\{2, 2.2\} = 2$$

To park

$$V[(1, 0)] = c = 4$$

$$V[(1, 1)] = \min\{1, 4\} = 1$$

To park



MDP vs MSP




CentraleSupélec

	MSP	MDP
Uncertainty	• Can model non-Markovian • Typically exogenous • Modeled as scenario tree	• Markovian • Naturally endogenous • Transition dynamics
Time horizon	• Typically finite	• Finite or infinite
State	• Often not explicitly defined	• Explicit state definition
Decision/ action	• Continuous or discrete • Naturally handle complex constraints	• Discrete (discretization or approximation if continuous) • Limited in handling constraints
Solution	• Decisions per scenario	• Policy mapping state $\rightarrow$ action
Complexity	• Curse of dimensionality • Approximate (e.g., decision rules)	• Curse of dimensionality • Approximate (e.g., value function approximation)



- How to find the optimal policy $\pi^* : S \rightarrow A$?



- Model-based methods:
 - 1) Value iteration
 - 2) Policy iteration
 - 3) Linear programming for MDPs
- Model-free (sampling): reinforcement learning



- The **optimal value** of a state s


$V^*(s)$ = expected utility starting in s
and acting optimally

- The **optimal value** of a Q-state (s, a)

$Q^*(s, a)$ = expected utility starting out
having taken action a from state s
and (thereafter) acting optimally

- The **optimal policy**

$\pi^*(s)$ = optimal action from state s







V^* and Q^*



CentraleSupélec



Noise = 0.2

Discount = 0.9

Living reward = 0




- **Fundamental operation:** compute the (expected max) **optimal value** of a state
 - Expected utility under **optimal** action
 - Average sum of (discounted) rewards

- **Recursive definition of values**
(Bellman optimality equations)

$$V^*(s) = \max_{a \in A} Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$






The Bellman equations



CentraleSupélec




How to be optimal:

Step 1: Take correct first action


Step 2: Keep being optimal



- Define $V_k(s)$ to be the **optimal value of s** if the decision horizon **ends in k more steps**
- How about infinite horizon (tree goes forever)?
 - idea: do a **depth-limited** computation, but with **increasing depths** until change is small
 - note: deep parts of the tree eventually don't matter if $\gamma < 1$



Decision horizon $K = 2$



$$V^*(s) = \max_{a \in A} \sum_{s' \in S} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



- **Idea:** repeatedly update an estimate of the optimal value function according to **Bellman optimality equation**
- **Procedure**
 - 1) Initialize $V_0(s) \leftarrow 0, \forall s \in S$: no time steps left means an expected reward sum of zero
 - 2) Update:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

3) Repeat until convergence

- Complexity of each iteration: $O(|S|^2|A|)$



Illustration of value iteration



CentraleSupélec

▲ 0.00	▲ 0.00	▲ 0.00	▲ 0.00
▲ 0.00		▲ 0.00	▲ 0.00
▲ 0.00	▲ 0.00	▲ 0.00	▲ 0.00
VALUES AFTER 0 ITERATIONS			

$$k = 0$$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

▲ 0.00	▲ 0.00	0.00 ▶	1.00
▲ 0.00		◀ 0.00	-1.00
▲ 0.00	▲ 0.00	▲ 0.00	0.00 ▼
VALUES AFTER 1 ITERATIONS			

$$k = 1$$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

▲ 0.00	0.00 ▶	0.72 ▶	1.00
▲ 0.00		▲ 0.00	-1.00
▲ 0.00	▲ 0.00	▲ 0.00	0.00 ▼
VALUES AFTER 2 ITERATIONS			

$$k = 2$$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

0.00 ▶	0.52 ▶	0.78 ▶	1.00
▲ 0.00		▲ 0.43	▼ -1.00
▲ 0.00	▲ 0.00	▲ 0.00	▼ 0.00
VALUES AFTER 3 ITERATIONS			

$$k = 3$$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

0.37 ▶	0.66 ▶	0.83 ▶	1.00
▲ 0.00		▲ 0.51	-1.00
▲ 0.00	0.00 ▶	▲ 0.31	◀ 0.00
VALUES AFTER 4 ITERATIONS			

$k = 4$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

0.51 ▶	0.72 ▶	0.84 ▶	1.00
▲ 0.27		▲ 0.55	-1.00
▲ 0.00	0.22 ▶	▲ 0.37	◀ 0.13
VALUES AFTER 5 ITERATIONS			

$$k = 5$$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

0.64 ▶	0.74 ▶	0.85 ▶	1.00
▲ 0.56		▲ 0.57	-1.00
▲ 0.48	◀ 0.41	▲ 0.47	◀ 0.27
VALUES AFTER 10 ITERATIONS			

$k = 10$

Noise = 0.2

Discount = 0.9

Living reward = 0



Illustration of value iteration



CentraleSupélec

0.64 ▶	0.74 ▶	0.85 ▶	1.00
▲ 0.57		▲ 0.57	-1.00
▲ 0.49	◀ 0.43	▲ 0.48	◀ 0.28
VALUES AFTER 100 ITERATIONS			

$k = 100$

Noise = 0.2

Discount = 0.9


Living reward = 0



Theorem: value iteration asymptotically converges to optimal value $V_k \rightarrow V^*$

Conceptual proof:

- **Case 1:** if the decision horizon is **finite** with T steps, then, V_T holds the actual untruncated values
- **Case 2:** if infinite horizon with $\gamma < 1$
- **Sketch:** for any state V_k and V_{k+1} can be viewed as depth $k+1$ expected-max results in **nearly identical** search trees
- The difference: on the bottom layer, V_{k+1} has actual rewards while V_k has zeros
- That last layer is at best all R_{MAX} , at worst R_{MIN} ; but everything is discounted by γ^k that far out, so



$$|V_{k+1} - V_k| \leq \gamma^k R_{MAX}$$



- Implementing the MSP formulation for the parking problem
- **Questions:** In the parking problem, is it possible to simplify the MSP model? How?



 yiping.fang@centralesupelec.fr